

HCT

In this notebook, we present a HCT survival prediction. In this notebook, compared to my previous starter notebooks we teach 5 new things:

- How to tranform `efs` and `efs_time` into single target with `KaplanMeierFitter` .
- How to train `GPU LightGBM model` with `KaplanMeierFitter` target
- How to train `XGBoost with Survivial:Cox loss`
- How to train `CatBoost with Survival:Cox loss`
- How to ensemble 5 models using `scipy.stats.rankdata()` .

Two Competition Approaches

In this competition, there are two ways to train a Survival Model:

- We can input both `efs` and `efs_time` and train a **model that supports** `survival loss like Cox` .
- Transform `efs` and `efs_time` into a single target proxy for `risk score` and train **any model** with `regression loss like MSE` .

In this notebook, we train 5 models. The first 3 models (XGBoost, CatBoost, LightGBM) use bullet point two. And the next 2 models (XGBoost Cox, CatBoost Cox) use bullet point one. Discussion about this notebook is [\[here\]\[4\]](#) and [\[here\]\[3\]](#).

Since this competition's metric is a ranking metric, we ensemble the 5 predictions by first converting each into ranks using `scipy.stats.rankdata()` . Afterward we created a weighted average from the ranks.

Have Fun! Enjoy!

Previous Notebooks

My previous starter notebooks are:

- XGBoost and CatBoost starter [\[here\]\[1\]](#)
- NN (MLP) starter [\[here\]\[2\]](#)

Pip Install Libraries for Metric

Since internet must be turned off for submission, we pip install from my other notebook [here](https://www.kaggle.com/code/cdeotte/pip-install-lifelines) (<https://www.kaggle.com/code/cdeotte/pip-install-lifelines>) where I downloaded the WHL files.

In [1]:

```
!pip install /kaggle/input/pip-install-lifelines/autograd-1.7.0-py3-none-any.whl
!pip install /kaggle/input/pip-install-lifelines/autograd-gamma-0.5.0.tar.gz
!pip install /kaggle/input/pip-install-lifelines/interface_meta-1.3.0-py3-none-any.whl
!pip install /kaggle/input/pip-install-lifelines/formulaic-1.0.2-py3-none-any.whl
!pip install /kaggle/input/pip-install-lifelines/lifelines-0.30.0-py3-none-any.whl
```

↕ Show hidden output

Load Train and Test

In [2]:

```

import numpy as np, pandas as pd
import matplotlib.pyplot as plt
pd.set_option('display.max_columns', 500)
pd.set_option('display.max_rows', 500)

test = pd.read_csv("/kaggle/input/equity-post-HCT-survival-predictions/
test.csv")
print("Test shape:", test.shape )

train = pd.read_csv("/kaggle/input/equity-post-HCT-survival-prediction
s/train.csv")
print("Train shape:", train.shape)
train.head()

```

Test shape: (3, 58)

Train shape: (28800, 60)

Out[2]:

	ID	dri_score	psych_disturb	cyto_score	diabetes	hla_match_c_high	hla_high_res_8
0	0	N/A - non-malignant indication	No	NaN	No	NaN	NaN
1	1	Intermediate	No	Intermediate	No	2.0	8.0
2	2	N/A - non-malignant indication	No	NaN	No	2.0	8.0
3	3	High	No	Intermediate	No	2.0	8.0
4	4	High	No	NaN	No	2.0	8.0

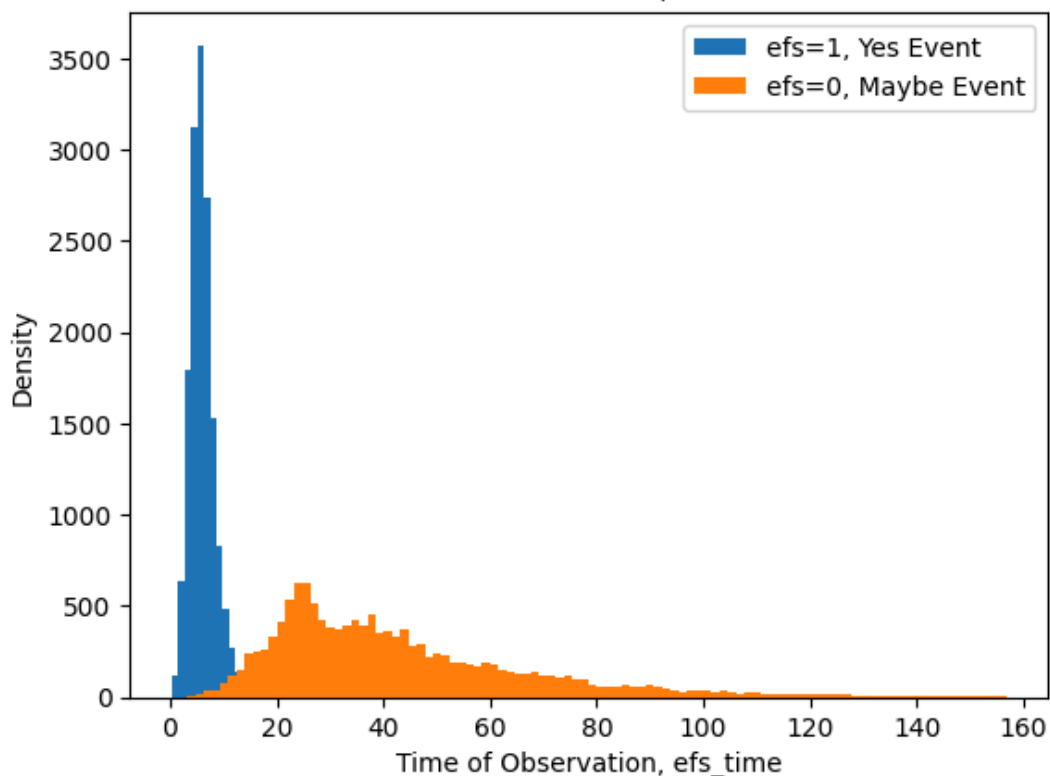
EDA on Train Targets

There are two train targets `efs` and `efs_time`. When `efs==1` we know patient **had an event** and we know time of event is `efs_time`. When `efs==0` we **do not know** if patient had an event or not, but we do know that patient was **without event for at least** `efs_time`.

In [3]:

```
plt.hist(train.loc[train.efs==1, "efs_time"], bins=100, label="efs=1, Yes Event")
plt.hist(train.loc[train.efs==0, "efs_time"], bins=100, label="efs=0, Maybe Event")
plt.xlabel("Time of Observation, efs_time")
plt.ylabel("Density")
plt.title("Times of Observation. Either time to event, or time observed without event.")
plt.legend()
plt.show()
```

Times of Observation. Either time to event, or time observed without event.



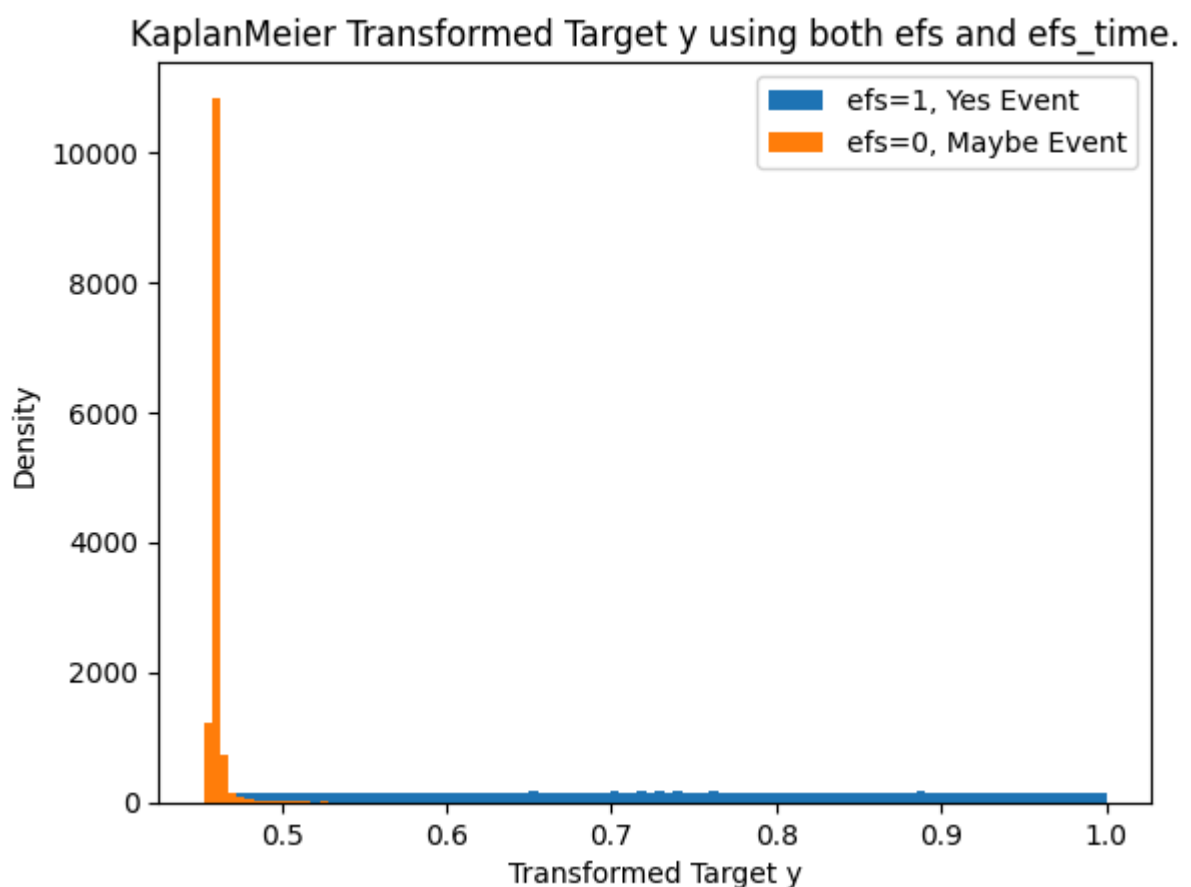
Transform Two Targets into One Target with KaplanMeier!

Both targets `efs` and `efs_time` provide useful information. We will transform these two targets into a single target to train our model with. In this competition we need to predict `risk score`. So we will create a target that mimics `risk score` to train our model. (Note this is only one out of many ways to transform two targets into one target. Considering experimenting on your own).

In [4]:

```
from lifelines import KaplanMeierFitter
def transform_survival_probability(df, time_col='efs_time', event_col='efs'):
    kmf = KaplanMeierFitter()
    kmf.fit(df[time_col], df[event_col])
    y = kmf.survival_function_at_times(df[time_col]).values
    return y
train["y"] = transform_survival_probability(train, time_col='efs_time',
event_col='efs')

plt.hist(train.loc[train.efs==1, "y"], bins=100, label="efs=1, Yes Event")
plt.hist(train.loc[train.efs==0, "y"], bins=100, label="efs=0, Maybe Event")
plt.xlabel("Transformed Target y")
plt.ylabel("Density")
plt.title("KaplanMeier Transformed Target y using both efs and efs_time.")
plt.legend()
plt.show()
```



Features

There are a total of 57 features. From these 35 are categorical and 22 are numerical. We will label encode the categorical features. Then our XGB and CAT model will accept these as categorical features and process them special internally. We leave the numerical feature NANs as NANs because GBDT (like XGB and CAT) can handle NAN and will use this information.

In [5]:

```
RMV = ["ID", "efs", "efs_time", "y"]
FEATURES = [c for c in train.columns if not c in RMV]
print(f"There are {len(FEATURES)} FEATURES: {FEATURES}")
```

```
There are 57 FEATURES: ['dri_score', 'psych_disturb', 'cyto_score',
'diabetes', 'hla_match_c_high', 'hla_high_res_8', 'tbi_status', 'ar
rhythmia', 'hla_low_res_6', 'graft_type', 'vent_hist', 'renal_issu
e', 'pulm_severe', 'prim_disease_hct', 'hla_high_res_6', 'cmv_statu
s', 'hla_high_res_10', 'hla_match_dqb1_high', 'tce_imm_match', 'hla
_nmdp_6', 'hla_match_c_low', 'rituximab', 'hla_match_drb1_low', 'hl
a_match_dqb1_low', 'prod_type', 'cyto_score_detail', 'conditioning_
intensity', 'ethnicity', 'year_hct', 'obesity', 'mrd_hct', 'in_vivo
_tcd', 'tce_match', 'hla_match_a_high', 'hepatic_severe', 'donor_ag
e', 'prior_tumor', 'hla_match_b_low', 'peptic_ulcer', 'age_at_hct',
'hla_match_a_low', 'gvhd_proph', 'rheum_issue', 'sex_match', 'hla_m
atch_b_high', 'race_group', 'comorbidity_score', 'karnofsky_score',
'hepatic_mild', 'tce_div_match', 'donor_related', 'melphalan_dose',
'hla_low_res_8', 'cardiac', 'hla_match_drb1_high', 'pulm_moderate',
'hla_low_res_10']
```

In [6]:

```
CATS = []
for c in FEATURES:
    if train[c].dtype=="object":
        CATS.append(c)
        train[c] = train[c].fillna("NAN")
        test[c] = test[c].fillna("NAN")
print(f"In these features, there are {len(CATS)} CATEGORICAL FEATURES: {CATS}")
```

In these features, there are 35 CATEGORICAL FEATURES: ['dri_score', 'psych_disturb', 'cyto_score', 'diabetes', 'tbi_status', 'arrhythmia', 'graft_type', 'vent_hist', 'renal_issue', 'pulm_severe', 'prim_disease_hct', 'cmv_status', 'tce_imm_match', 'rituximab', 'prod_type', 'cyto_score_detail', 'conditioning_intensity', 'ethnicity', 'obesity', 'mrd_hct', 'in_vivo_tcd', 'tce_match', 'hepatic_severe', 'prior_tumor', 'peptic_ulcer', 'gvhd_proph', 'rheum_issue', 'sex_match', 'race_group', 'hepatic_mild', 'tce_div_match', 'donor_related', 'melphalan_dose', 'cardiac', 'pulm_moderate']

In [7]:

```

combined = pd.concat([train,test],axis=0,ignore_index=True)
#print("Combined data shape:", combined.shape )

# LABEL ENCODE CATEGORICAL FEATURES
print("We LABEL ENCODE the CATEGORICAL FEATURES: ",end="")
for c in FEATURES:

    # LABEL ENCODE CATEGORICAL AND CONVERT TO INT32 CATEGORY
    if c in CATS:
        print(f"{c}, ",end="")
        combined[c],_ = combined[c].factorize()
        combined[c] -= combined[c].min()
        combined[c] = combined[c].astype("int32")
        combined[c] = combined[c].astype("category")

    # REDUCE PRECISION OF NUMERICAL TO 32BIT TO SAVE MEMORY
    else:
        if combined[c].dtype=="float64":
            combined[c] = combined[c].astype("float32")
        if combined[c].dtype=="int64":
            combined[c] = combined[c].astype("int32")

train = combined.iloc[:len(train)].copy()
test = combined.iloc[len(train):].reset_index(drop=True).copy()

```

We LABEL ENCODE the CATEGORICAL FEATURES: dri_score, psych_disturb, cyto_score, diabetes, tbi_status, arrhythmia, graft_type, vent_his t, renal_issue, pulm_severe, prim_disease_hct, cmv_status, tce_imm match, rituximab, prod_type, cyto_score_detail, conditioning_intens ity, ethnicity, obesity, mrd_hct, in_vivo_tcd, tce_match, hepatic_s evere, prior_tumor, peptic_ulcer, gvhd_proph, rheum_issue, sex_matc h, race_group, hepatic_mild, tce_div_match, donor_related, melphala n_dose, cardiac, pulm_moderate,

XGBoost with KaplanMeier

We train XGBoost model for 10 folds and achieve **CV 0.674!**

In [8]:

```
from sklearn.model_selection import KFold
from xgboost import XGBRegressor, XGBClassifier
import xgboost as xgb
print("Using XGBoost version", xgb.__version__)
```

Using XGBoost version 2.0.3

In [9]:

```
%%time
FOLDS = 10
kf = KFold(n_splits=FOLDS, shuffle=True, random_state=42)

oof_xgb = np.zeros(len(train))
pred_xgb = np.zeros(len(test))

for i, (train_index, test_index) in enumerate(kf.split(train)):

    print("#"*25)
    print(f"### Fold {i+1}")
    print("#"*25)

    x_train = train.loc[train_index, FEATURES].copy()
    y_train = train.loc[train_index, "y"]
    x_valid = train.loc[test_index, FEATURES].copy()
    y_valid = train.loc[test_index, "y"]
    x_test = test[FEATURES].copy()

    model_xgb = XGBRegressor(
        device="cuda",
        max_depth=3,
        colsample_bytree=0.5,
        subsample=0.8,
        n_estimators=2000,
        learning_rate=0.02,
        enable_categorical=True,
        min_child_weight=80,
        #early_stopping_rounds=25,
    )
    model_xgb.fit(
        x_train, y_train,
        eval_set=[(x_valid, y_valid)],
        verbose=500
    )

    # INFER OOF
    oof_xgb[test_index] = model_xgb.predict(x_valid)
    # INFER TEST
    pred_xgb += model_xgb.predict(x_test)
```

```
# COMPUTE AVERAGE TEST PREDS  
pred_xgb /= FOLDS
```

```
#####  
### Fold 1  
#####  
[0]      validation_0-rmse:0.17773  
[500]    validation_0-rmse:0.15956  
[1000]   validation_0-rmse:0.15746  
[1500]   validation_0-rmse:0.15650  
[1999]   validation_0-rmse:0.15605  
#####  
### Fold 2  
#####  
[0]      validation_0-rmse:0.17350
```

/opt/conda/lib/python3.10/site-packages/xgboost/core.py:160: UserWarning: [05:40:06] WARNING: /workspace/src/common/error_msg.cc:58: Falling back to prediction using DMatrix due to mismatched devices. This might lead to higher memory usage and slower performance. XGBoost is running on: cuda:0, while the input data is on: cpu.

Potential solutions:

- Use a data structure that matches the device ordinal in the booster.
- Set the device for booster before call to inplace_predict.

This warning will only be shown once.

```
warnings.warn(smsg, UserWarning)
```

```
[500] validation_0-rmse:0.15572
[1000] validation_0-rmse:0.15422
[1500] validation_0-rmse:0.15356
[1999] validation_0-rmse:0.15312
#####
### Fold 3
#####
[0] validation_0-rmse:0.17724
[500] validation_0-rmse:0.15800
[1000] validation_0-rmse:0.15612
[1500] validation_0-rmse:0.15538
[1999] validation_0-rmse:0.15494
#####
### Fold 4
#####
[0] validation_0-rmse:0.17923
[500] validation_0-rmse:0.16024
[1000] validation_0-rmse:0.15808
[1500] validation_0-rmse:0.15713
[1999] validation_0-rmse:0.15668
#####
### Fold 5
#####
[0] validation_0-rmse:0.17368
[500] validation_0-rmse:0.15737
[1000] validation_0-rmse:0.15550
[1500] validation_0-rmse:0.15474
[1999] validation_0-rmse:0.15431
#####
### Fold 6
#####
[0] validation_0-rmse:0.17748
[500] validation_0-rmse:0.15964
[1000] validation_0-rmse:0.15802
[1500] validation_0-rmse:0.15738
[1999] validation_0-rmse:0.15698
#####
### Fold 7
#####
[0] validation_0-rmse:0.17846
[500] validation_0-rmse:0.16149
[1000] validation_0-rmse:0.15980
[1500] validation_0-rmse:0.15907
```

```
[1999] validation_0-rmse:0.15873
#####
### Fold 8
#####
[0] validation_0-rmse:0.17457
[500] validation_0-rmse:0.15782
[1000] validation_0-rmse:0.15580
[1500] validation_0-rmse:0.15495
[1999] validation_0-rmse:0.15453
#####
### Fold 9
#####
[0] validation_0-rmse:0.17648
[500] validation_0-rmse:0.16014
[1000] validation_0-rmse:0.15847
[1500] validation_0-rmse:0.15762
[1999] validation_0-rmse:0.15715
#####
### Fold 10
#####
[0] validation_0-rmse:0.17527
[500] validation_0-rmse:0.15816
[1000] validation_0-rmse:0.15628
[1500] validation_0-rmse:0.15545
[1999] validation_0-rmse:0.15506
CPU times: user 57.6 s, sys: 383 ms, total: 58 s
Wall time: 52.7 s
```

In [10]:


```
from metric import score

y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = oof_xgb
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for XGBoost KaplanMeier =", m)
```

/kaggle/usr/lib/efs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for XGBoost KaplanMeier = 0.6737940261928012

[Show hidden code](#)

CatBoost with KaplanMeier

We train CatBoost model for 10 folds and achieve **CV 0.674!**

In [12]:

```
from catboost import CatBoostRegressor, CatBoostClassifier
import catboost as cb
print("Using CatBoost version",cb.__version__)
```

Using CatBoost version 1.2.7

In [13]:

```
%%time
FOLDS = 10
kf = KFold(n_splits=FOLDS, shuffle=True, random_state=42)

oof_cat = np.zeros(len(train))
pred_cat = np.zeros(len(test))

for i, (train_index, test_index) in enumerate(kf.split(train)):

    print("#"*25)
    print(f"### Fold {i+1}")
    print("#"*25)

    x_train = train.loc[train_index, FEATURES].copy()
    y_train = train.loc[train_index, "y"]
    x_valid = train.loc[test_index, FEATURES].copy()
    y_valid = train.loc[test_index, "y"]
    x_test = test[FEATURES].copy()

    model_cat = CatBoostRegressor(
        task_type="GPU",
        learning_rate=0.1,
        grow_policy='Lossguide',
        #early_stopping_rounds=25,
    )
    model_cat.fit(x_train, y_train,
                  eval_set=(x_valid, y_valid),
                  cat_features=CATS,
                  verbose=250)

    # INFER OOF
    oof_cat[test_index] = model_cat.predict(x_valid)
    # INFER TEST
    pred_cat += model_cat.predict(x_test)

# COMPUTE AVERAGE TEST PREDS
pred_cat /= FOLDS
```

#####

Fold 1

#####

0: learn: 0.1743993 test: 0.1760769 best: 0.1760769 (0)
total: 100ms remaining: 1m 40s
250: learn: 0.1444360 test: 0.1574516 best: 0.1573758 (24
3) total: 5.21s remaining: 15.5s
500: learn: 0.1364002 test: 0.1570436 best: 0.1569415 (48
6) total: 10.5s remaining: 10.4s
750: learn: 0.1295978 test: 0.1567639 best: 0.1566160 (69
4) total: 15.5s remaining: 5.15s
999: learn: 0.1238415 test: 0.1566584 best: 0.1566160 (69
4) total: 20s remaining: 0us

bestTest = 0.1566160111

bestIteration = 694

Shrink model to first 695 iterations.

#####

Fold 2

#####

0: learn: 0.1748674 test: 0.1717315 best: 0.1717315 (0)
total: 24ms remaining: 24s
250: learn: 0.1445834 test: 0.1532805 best: 0.1532805 (25
0) total: 4.67s remaining: 13.9s
500: learn: 0.1361701 test: 0.1527256 best: 0.1527189 (48
9) total: 9.01s remaining: 8.98s
750: learn: 0.1297626 test: 0.1529458 best: 0.1526560 (55
5) total: 13.7s remaining: 4.54s
999: learn: 0.1241974 test: 0.1531554 best: 0.1526560 (55
5) total: 20s remaining: 0us

bestTest = 0.1526560481

bestIteration = 555

Shrink model to first 556 iterations.

#####

Fold 3

#####

0: learn: 0.1745178 test: 0.1754648 best: 0.1754648 (0)
total: 20.1ms remaining: 20.1s
250: learn: 0.1440720 test: 0.1555157 best: 0.1555074 (24
6) total: 4.54s remaining: 13.6s
500: learn: 0.1358665 test: 0.1552825 best: 0.1552369 (43
0) total: 9.72s remaining: 9.68s
750: learn: 0.1295579 test: 0.1552502 best: 0.1550891 (60
1) total: 15s remaining: 4.99s

```
999:   learn: 0.1240175          test: 0.1550098 best: 0.1549959 (95
7)     total: 20.4s    remaining: 0us
bestTest = 0.1549959338
bestIteration = 957
Shrink model to first 958 iterations.
#####
### Fold 4
#####
0:      learn: 0.1742386          test: 0.1775053 best: 0.1775053 (0)
total: 22ms    remaining: 22s
250:    learn: 0.1443504          test: 0.1572211 best: 0.1571940 (24
9)     total: 4.67s    remaining: 13.9s
500:    learn: 0.1362868          test: 0.1564467 best: 0.1564252 (46
5)     total: 10.1s    remaining: 10.1s
750:    learn: 0.1297536          test: 0.1560188 best: 0.1560188 (75
0)     total: 14.5s    remaining: 4.82s
999:    learn: 0.1242869          test: 0.1560844 best: 0.1559439 (89
7)     total: 19s      remaining: 0us
bestTest = 0.1559438801
bestIteration = 897
Shrink model to first 898 iterations.
#####
### Fold 5
#####
0:      learn: 0.1748209          test: 0.1720765 best: 0.1720765 (0)
total: 21.6ms   remaining: 21.6s
250:    learn: 0.1450010          test: 0.1545487 best: 0.1545283 (24
9)     total: 4.06s    remaining: 12.1s
500:    learn: 0.1367038          test: 0.1543137 best: 0.1541771 (37
6)     total: 9.21s    remaining: 9.18s
750:    learn: 0.1301713          test: 0.1545665 best: 0.1541771 (37
6)     total: 14.4s    remaining: 4.78s
999:    learn: 0.1246478          test: 0.1548056 best: 0.1541771 (37
6)     total: 19.5s    remaining: 0us
bestTest = 0.1541770502
bestIteration = 376
Shrink model to first 377 iterations.
#####
### Fold 6
#####
0:      learn: 0.1744351          test: 0.1757309 best: 0.1757309 (0)
total: 18.7ms   remaining: 18.7s
250:    learn: 0.1449407          test: 0.1572496 best: 0.1572496 (25
0)     total: 4.04s    remaining: 12s
```

```
500:   learn: 0.1371375          test: 0.1571872 best: 0.1570450 (37
9)     total: 8.59s    remaining: 8.56s
750:   learn: 0.1308869          test: 0.1573460 best: 0.1570450 (37
9)     total: 13.9s    remaining: 4.59s
999:   learn: 0.1252236          test: 0.1579734 best: 0.1570450 (37
9)     total: 19.4s    remaining: 0us
```

```
bestTest = 0.1570450034
```

```
bestIteration = 379
```

```
Shrink model to first 380 iterations.
```

```
#####
```

```
### Fold 7
```

```
#####
```

```
0:      learn: 0.1743394          test: 0.1767087 best: 0.1767087 (0)
total: 21.4ms    remaining: 21.4s
250:    learn: 0.1446544          test: 0.1592181 best: 0.1592041 (24
7)      total: 4.89s    remaining: 14.6s
500:    learn: 0.1364501          test: 0.1588536 best: 0.1587487 (46
2)      total: 10s      remaining: 10s
750:    learn: 0.1299135          test: 0.1592445 best: 0.1587487 (46
2)      total: 15s      remaining: 4.98s
999:    learn: 0.1244221          test: 0.1593463 best: 0.1587487 (46
2)      total: 19.5s    remaining: 0us
```

```
bestTest = 0.1587486709
```

```
bestIteration = 462
```

```
Shrink model to first 463 iterations.
```

```
#####
```

```
### Fold 8
```

```
#####
```

```
0:      learn: 0.1746938          test: 0.1728086 best: 0.1728086 (0)
total: 21.4ms    remaining: 21.4s
250:    learn: 0.1442486          test: 0.1553799 best: 0.1553799 (25
0)      total: 4.78s    remaining: 14.3s
500:    learn: 0.1361261          test: 0.1548325 best: 0.1548117 (45
4)      total: 9.83s    remaining: 9.79s
750:    learn: 0.1293713          test: 0.1544145 best: 0.1543675 (73
5)      total: 15.2s    remaining: 5.03s
999:    learn: 0.1236059          test: 0.1545633 best: 0.1543675 (73
5)      total: 20.6s    remaining: 0us
```

```
bestTest = 0.154367457
```

```
bestIteration = 735
```

```
Shrink model to first 736 iterations.
```

```
#####
```

```
### Fold 9
```

```
#####
```

```
0:      learn: 0.1744840      test: 0.1748802 best: 0.1748802 (0)
total: 21.9ms   remaining: 21.9s
250:    learn: 0.1442797      test: 0.1583441 best: 0.1583423 (24
5)      total: 5.01s   remaining: 15s
500:    learn: 0.1361857      test: 0.1577813 best: 0.1577649 (49
6)      total: 10.3s   remaining: 10.2s
750:    learn: 0.1297267      test: 0.1573862 best: 0.1573524 (74
7)      total: 15.5s   remaining: 5.15s
999:    learn: 0.1241866      test: 0.1573732 best: 0.1573524 (74
7)      total: 20.7s   remaining: 0us
bestTest = 0.1573523566
bestIteration = 747
Shrink model to first 748 iterations.
#####
### Fold 10
#####
0:      learn: 0.1746737      test: 0.1735036 best: 0.1735036 (0)
total: 21.7ms   remaining: 21.7s
250:    learn: 0.1453132      test: 0.1558658 best: 0.1558658 (25
0)      total: 4.98s   remaining: 14.9s
500:    learn: 0.1373193      test: 0.1556390 best: 0.1555556 (35
7)      total: 10.1s   remaining: 10s
750:    learn: 0.1311796      test: 0.1558097 best: 0.1555293 (68
1)      total: 15.6s   remaining: 5.16s
999:    learn: 0.1258469      test: 0.1561605 best: 0.1555293 (68
1)      total: 21s     remaining: 0us
bestTest = 0.1555292758
bestIteration = 681
Shrink model to first 682 iterations.
CPU times: user 8min 27s, sys: 2min 34s, total: 11min 2s
Wall time: 3min 31s
```

In [14]:


```
y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = oof_cat
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for CatBoost KaplanMeier =", m)
```

/kaggle/usr/lib/efs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for CatBoost KaplanMeier = 0.6740795777257533

⌵ Show hidden code



LightGBM with KaplanMeier

We train LightGBM model for 10 folds and achieve **CV 0.6725!**

In [16]:

```
from lightgbm import LGBMRegressor
import lightgbm as lgb
print("Using LightGBM version", lgb.__version__)
```

Using LightGBM version 4.2.0

In [17]:

```
FOLDS = 10
kf = KFold(n_splits=FOLDS, shuffle=True, random_state=42)

oof_lgb = np.zeros(len(train))
pred_lgb = np.zeros(len(test))

for i, (train_index, test_index) in enumerate(kf.split(train)):

    print("#"*25)
    print(f"### Fold {i+1}")
    print("#"*25)

    x_train = train.loc[train_index, FEATURES].copy()
    y_train = train.loc[train_index, "y"]
    x_valid = train.loc[test_index, FEATURES].copy()
    y_valid = train.loc[test_index, "y"]
    x_test = test[FEATURES].copy()

    model_lgb = LGBMRegressor(
        device="gpu",
        max_depth=3,
        colsample_bytree=0.4,
        #subsample=0.9,
        n_estimators=2500,
        learning_rate=0.02,
        objective="regression",
        verbose=-1,
        #early_stopping_rounds=25,
    )
    model_lgb.fit(
        x_train, y_train,
        eval_set=[(x_valid, y_valid)],
    )

    # INFER OOF
    oof_lgb[test_index] = model_lgb.predict(x_valid)
    # INFER TEST
    pred_lgb += model_lgb.predict(x_test)

# COMPUTE AVERAGE TEST PREDS
pred_lgb /= FOLDS
```

[illegible]

```
#####  
### Fold 2  
#####  
#####  
### Fold 3  
#####  
#####  
### Fold 4  
#####  
#####  
### Fold 5  
#####  
#####  
### Fold 6  
#####  
#####  
### Fold 7  
#####  
#####  
### Fold 8  
#####  
#####  
### Fold 9  
#####  
#####  
### Fold 10  
#####
```

In [18]:


```
y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = oof_lgb
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for LightGBM KaplanMeier =", m)
```

/kaggle/usr/lib/efs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for LightGBM KaplanMeier = 0.6725169670618927

↕ Show hidden code



XGBoost with Survival:Cox

We train XGBoost using Survival:Cox loss for 10 folds and achieve **CV=672!**

In [20]:

```
# SURVIVAL COX NEEDS THIS TARGET (TO DIGEST EFS AND EFS_TIME)  
train["efs_time2"] = train.efs_time.copy()  
train.loc[train.efs==0, "efs_time2"] *= -1
```


In [21]:

```
FOLDS = 10
kf = KFold(n_splits=FOLDS, shuffle=True, random_state=42)

oof_xgb_cox = np.zeros(len(train))
pred_xgb_cox = np.zeros(len(test))

for i, (train_index, test_index) in enumerate(kf.split(train)):

    print("#"*25)
    print(f"### Fold {i+1}")
    print("#"*25)

    x_train = train.loc[train_index, FEATURES].copy()
    y_train = train.loc[train_index, "efs_time2"]
    x_valid = train.loc[test_index, FEATURES].copy()
    y_valid = train.loc[test_index, "efs_time2"]
    x_test = test[FEATURES].copy()

    model_xgb_cox = XGBRegressor(
        device="cuda",
        max_depth=3,
        colsample_bytree=0.5,
        subsample=0.8,
        n_estimators=2000,
        learning_rate=0.02,
        enable_categorical=True,
        min_child_weight=80,
        objective='survival:cox',
        eval_metric='cox-nloglik',
    )
    model_xgb_cox.fit(
        x_train, y_train,
        eval_set=[(x_valid, y_valid)],
        verbose=500
    )

    # INFER OOF
    oof_xgb_cox[test_index] = model_xgb_cox.predict(x_valid)
    # INFER TEST
    pred_xgb_cox += model_xgb_cox.predict(x_test)
```

```
# COMPUTE AVERAGE TEST PREDS  
pred_xgb_cox /= FOLDS
```

```
#####  
### Fold 1  
#####  
[0] validation_0-cox-nloglik:7.62402  
[500] validation_0-cox-nloglik:7.43513  
[1000] validation_0-cox-nloglik:7.41916  
[1500] validation_0-cox-nloglik:7.41221  
[1999] validation_0-cox-nloglik:7.41085  
#####  
### Fold 2  
#####  
[0] validation_0-cox-nloglik:7.61704  
[500] validation_0-cox-nloglik:7.41060  
[1000] validation_0-cox-nloglik:7.39630  
[1500] validation_0-cox-nloglik:7.39059  
[1999] validation_0-cox-nloglik:7.38769  
#####  
### Fold 3  
#####  
[0] validation_0-cox-nloglik:7.60952  
[500] validation_0-cox-nloglik:7.40543  
[1000] validation_0-cox-nloglik:7.39056  
[1500] validation_0-cox-nloglik:7.38663  
[1999] validation_0-cox-nloglik:7.38550  
#####  
### Fold 4  
#####  
[0] validation_0-cox-nloglik:7.60515  
[500] validation_0-cox-nloglik:7.41071  
[1000] validation_0-cox-nloglik:7.40056  
[1500] validation_0-cox-nloglik:7.39666  
[1999] validation_0-cox-nloglik:7.39720  
#####  
### Fold 5  
#####  
[0] validation_0-cox-nloglik:7.62895  
[500] validation_0-cox-nloglik:7.42383  
[1000] validation_0-cox-nloglik:7.40763  
[1500] validation_0-cox-nloglik:7.40042  
[1999] validation_0-cox-nloglik:7.39660  
#####  
### Fold 6  
#####
```

```
[0] validation_0-cox-nloglik:7.61275
[500] validation_0-cox-nloglik:7.40221
[1000] validation_0-cox-nloglik:7.39105
[1500] validation_0-cox-nloglik:7.38766
[1999] validation_0-cox-nloglik:7.38749
#####
### Fold 7
#####
[0] validation_0-cox-nloglik:7.62944
[500] validation_0-cox-nloglik:7.44143
[1000] validation_0-cox-nloglik:7.42778
[1500] validation_0-cox-nloglik:7.42296
[1999] validation_0-cox-nloglik:7.42149
#####
### Fold 8
#####
[0] validation_0-cox-nloglik:7.61662
[500] validation_0-cox-nloglik:7.44155
[1000] validation_0-cox-nloglik:7.42844
[1500] validation_0-cox-nloglik:7.42232
[1999] validation_0-cox-nloglik:7.41951
#####
### Fold 9
#####
[0] validation_0-cox-nloglik:7.61684
[500] validation_0-cox-nloglik:7.44634
[1000] validation_0-cox-nloglik:7.43420
[1500] validation_0-cox-nloglik:7.42970
[1999] validation_0-cox-nloglik:7.42799
#####
### Fold 10
#####
[0] validation_0-cox-nloglik:7.61647
[500] validation_0-cox-nloglik:7.43413
[1000] validation_0-cox-nloglik:7.42128
[1500] validation_0-cox-nloglik:7.41845
[1999] validation_0-cox-nloglik:7.41635
```

In [22]:


```
y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = oof_xgb_cox
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for XGBoost Survival:Cox =", m)
```

/kaggle/usr/lib/efs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for XGBoost Survival:Cox = 0.6722446470179296

⌵ Show hidden code



CatBoost with Survival:Cox

We train CatBoost using Survival:Cox loss for 10 folds and achieve **CV=671!**

In [24]:

```
FOLDS = 10
kf = KFold(n_splits=FOLDS, shuffle=True, random_state=42)

oof_cat_cox = np.zeros(len(train))
pred_cat_cox = np.zeros(len(test))

for i, (train_index, test_index) in enumerate(kf.split(train)):

    print("#"*25)
    print(f"### Fold {i+1}")
    print("#"*25)

    x_train = train.loc[train_index, FEATURES].copy()
    y_train = train.loc[train_index, "efs_time2"]
    x_valid = train.loc[test_index, FEATURES].copy()
    y_valid = train.loc[test_index, "efs_time2"]
    x_test = test[FEATURES].copy()

    model_cat_cox = CatBoostRegressor(
        loss_function="Cox",
        #task_type="GPU",
        iterations=400,
        learning_rate=0.1,
        grow_policy='Lossguide',
        use_best_model=False,
    )
    model_cat_cox.fit(x_train, y_train,
                      eval_set=(x_valid, y_valid),
                      cat_features=CATS,
                      verbose=100)

    # INFER OOF
    oof_cat_cox[test_index] = model_cat_cox.predict(x_valid)
    # INFER TEST
    pred_cat_cox += model_cat_cox.predict(x_test)

# COMPUTE AVERAGE TEST PREDS
pred_cat_cox /= FOLDS
```

#####

Fold 1

#####

0:	learn: -137204.2010418	test: -11625.0126498	best: -1162
5.0126498 (0)	total: 70.9ms	remaining: 28.3s	
100:	learn: -134245.0940003	test: -11368.0935757	best: -1136
7.7720241 (99)	total: 5.81s	remaining: 17.2s	
200:	learn: -133569.4247640	test: -11357.0053940	best: -1135
6.8330165 (182)	total: 11.6s	remaining: 11.5s	
300:	learn: -133095.7842781	test: -11351.1819262	best: -1135
1.0222775 (299)	total: 17.4s	remaining: 5.71s	
399:	learn: -132763.5913301	test: -11349.4816640	best: -1134
9.4142821 (327)	total: 22.9s	remaining: 0us	

bestTest = -11349.41428

bestIteration = 327

#####

Fold 2

#####

0:	learn: -137014.2912101	test: -11772.8856048	best: -1177
2.8856048 (0)	total: 63.7ms	remaining: 25.4s	
100:	learn: -134091.3022715	test: -11485.4489792	best: -1148
5.3225232 (99)	total: 6.23s	remaining: 18.4s	
200:	learn: -133312.7852628	test: -11460.6629034	best: -1146
0.6629034 (200)	total: 12.1s	remaining: 11.9s	
300:	learn: -132843.8300906	test: -11453.5101666	best: -1145
3.1395642 (286)	total: 17.9s	remaining: 5.87s	
399:	learn: -132444.2041710	test: -11451.6650578	best: -1145
1.1640114 (386)	total: 23.6s	remaining: 0us	

bestTest = -11451.16401

bestIteration = 386

#####

Fold 3

#####

0:	learn: -136740.2719659	test: -11983.0664595	best: -1198
3.0664595 (0)	total: 64.8ms	remaining: 25.9s	
100:	learn: -133765.3366558	test: -11689.7400344	best: -1168
9.7400344 (100)	total: 5.81s	remaining: 17.2s	
200:	learn: -133055.1524830	test: -11675.0143694	best: -1167
4.4228636 (194)	total: 12.1s	remaining: 11.9s	


```
300:   learn: -132628.9478783 test: -11670.8603836 best: -1167
0.7024139 (293) total: 17.9s remaining: 5.88s
399:   learn: -132318.8285745 test: -11674.4124251 best: -1167
0.3801276 (317) total: 23.5s remaining: 0us
```

```
bestTest = -11670.38013
```

```
bestIteration = 317
```

```
#####
```

```
### Fold 4
```

```
#####
```

```
0:   learn: -136474.7243316 test: -12180.0536823 best: -1218
0.0536823 (0) total: 64.1ms remaining: 25.6s
100: learn: -133463.0770737 test: -11892.2690720 best: -1189
2.2690720 (100) total: 5.81s remaining: 17.2s
200: learn: -132783.0162317 test: -11878.1443791 best: -1187
7.4964925 (197) total: 11.6s remaining: 11.5s
300: learn: -132368.6648703 test: -11875.1126818 best: -1187
4.8392905 (290) total: 17.4s remaining: 5.72s
399: learn: -131959.4648801 test: -11873.9173657 best: -1187
3.8885020 (398) total: 23.4s remaining: 0us
```

```
bestTest = -11873.8885
```

```
bestIteration = 398
```

```
#####
```

```
### Fold 5
```

```
#####
```

```
0:   learn: -137321.8175168 test: -11539.7868480 best: -1153
9.7868480 (0) total: 68.7ms remaining: 27.4s
100: learn: -134353.3215180 test: -11253.3822723 best: -1125
3.3822723 (100) total: 5.8s remaining: 17.2s
200: learn: -133623.9197023 test: -11235.6006972 best: -1123
5.1441845 (198) total: 11.6s remaining: 11.5s
300: learn: -133129.5479754 test: -11236.6138716 best: -1123
3.6543759 (259) total: 17.4s remaining: 5.72s
399: learn: -132739.6146843 test: -11237.5755754 best: -1123
3.6543759 (259) total: 23.1s remaining: 0us
```

```
bestTest = -11233.65438
```

```
bestIteration = 259
```

```
#####
```

```
### Fold 6
```

#####

0:	learn: -136830.6351496	test: -11908.3484761	best: -1190
8.3484761 (0)	total: 64ms	remaining: 25.5s	
100:	learn: -133900.6412401	test: -11611.3469780	best: -1161
1.3469780 (100)	total: 6.18s	remaining: 18.3s	
200:	learn: -133046.9692152	test: -11603.0241538	best: -1160
1.2102419 (171)	total: 12s	remaining: 11.9s	
300:	learn: -132618.4127912	test: -11601.3943921	best: -1159
9.6831955 (279)	total: 17.8s	remaining: 5.87s	
399:	learn: -132280.1838119	test: -11603.5361790	best: -1159
9.6831955 (279)	total: 23.6s	remaining: 0us	

bestTest = -11599.6832

bestIteration = 279

#####

Fold 7

#####

0:	learn: -137331.1086384	test: -11527.8348135	best: -1152
7.8348135 (0)	total: 64.4ms	remaining: 25.7s	
100:	learn: -134301.3515021	test: -11269.2842488	best: -1126
9.2842488 (100)	total: 5.85s	remaining: 17.3s	
200:	learn: -133538.0266711	test: -11258.8797675	best: -1125
8.6810837 (195)	total: 11.8s	remaining: 11.7s	
300:	learn: -133115.9160399	test: -11261.4263600	best: -1125
8.6810837 (195)	total: 18s	remaining: 5.91s	
399:	learn: -132757.7938084	test: -11266.3943387	best: -1125
8.6810837 (195)	total: 23.6s	remaining: 0us	

bestTest = -11258.68108

bestIteration = 195

#####

Fold 8

#####

0:	learn: -136894.4434350	test: -11865.5004882	best: -1186
5.5004882 (0)	total: 63.8ms	remaining: 25.5s	
100:	learn: -133874.3737942	test: -11612.4429632	best: -1161
2.4429632 (100)	total: 5.77s	remaining: 17.1s	
200:	learn: -133157.3561881	test: -11599.3268855	best: -1159
8.9962380 (191)	total: 11.6s	remaining: 11.5s	
300:	learn: -132817.2931471	test: -11599.7885483	best: -1159
8.5509036 (220)	total: 17.3s	remaining: 5.7s	
399:	learn: -132509.2524231	test: -11598.2225128	best: -1159

5.7294350 (338) total: 23.4s remaining: 0us

bestTest = -11595.72943

bestIteration = 338

#####

Fold 9

#####

0:	learn: -136897.9544760	test: -11860.2823079	best: -1186
0.2823079 (0)	total: 62.4ms	remaining: 24.9s	
100:	learn: -133935.9459575	test: -11622.3451615	best: -1162
2.3451615 (100)	total: 5.8s	remaining: 17.2s	
200:	learn: -133124.8259221	test: -11615.2248965	best: -1161
3.1851477 (149)	total: 11.6s	remaining: 11.5s	
300:	learn: -132651.7234484	test: -11615.8087751	best: -1161
3.0622011 (252)	total: 17.5s	remaining: 5.75s	
399:	learn: -132356.1183940	test: -11616.0147572	best: -1161
3.0622011 (252)	total: 23.2s	remaining: 0us	

bestTest = -11613.0622

bestIteration = 252

#####

Fold 10

#####

0:	learn: -136968.8520725	test: -11803.6923015	best: -1180
3.6923015 (0)	total: 61.9ms	remaining: 24.7s	
100:	learn: -133987.6775754	test: -11539.6778644	best: -1153
9.6778644 (100)	total: 6.02s	remaining: 17.8s	
200:	learn: -133261.2056712	test: -11531.2596141	best: -1152
9.6792268 (160)	total: 12.1s	remaining: 12s	
300:	learn: -132741.7177251	test: -11527.6029744	best: -1152
7.5338030 (297)	total: 17.9s	remaining: 5.89s	
399:	learn: -132387.7444014	test: -11529.0931774	best: -1152
6.9063086 (303)	total: 23.6s	remaining: 0us	

bestTest = -11526.90631

bestIteration = 303

In [25]:


```
y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = oof_cat_cox
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for CatBoost Survival:Cox =", m)
```

/kaggle/usr/lib/efs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for CatBoost Survival:Cox = 0.6707201406861238

⌵ Show hidden code



Ensemble CAT and XGB and LGB

We ensemble our XGBoost, CatBoost, LightGBM, XGBoost Cox, and CatBoost Cox using

`scipy.stats.rankdata()` and achieve an amazing **CV=0.681** Wow!

In [27]:

```
from scipy.stats import rankdata

y_true = train[["ID", "efs", "efs_time", "race_group"]].copy()
y_pred = train[["ID"]].copy()
y_pred["prediction"] = rankdata(oof_xgb) + rankdata(oof_cat) + rankdata(
    oof_lgb) \
    + rankdata(oof_xgb_cox) + rankdata(oof_cat_cox)
m = score(y_true.copy(), y_pred.copy(), "ID")
print(f"\nOverall CV for Ensemble =", m)
```

/kaggle/usr/lib/eeefs-concordance-index/metric.py:59: FutureWarning:
The default of observed=False is deprecated and will be changed to
True in a future version of pandas. Pass observed=False to retain c
urrent behavior or observed=True to adopt the future default and si
lence this warning.

```
merged_df_race_dict = dict(merged_df.groupby(['race_group']).grou  
ps)
```

Overall CV for Ensemble = 0.680877349730958

Create Submission CSV

In [28]:

```
sub = pd.read_csv("/kaggle/input/equity-post-HCT-survival-predictions/s
ample_submission.csv")
sub.prediction = rankdata(pred_xgb) + rankdata(pred_cat) + rankdata(pre
d_lgb)\
                    + rankdata(pred_xgb_cox) + rankdata(pred_cat_cox)
sub.to_csv("submission.csv", index=False)
print("Sub shape:", sub.shape)
sub.head()
```

Sub shape: (3, 2)

Out[28]:

	ID	prediction
0	28800	10.0
1	28801	15.0
2	28802	5.0

In [29]:

```
print(model_xgb)
print(model_xgb_cox)
print(model_cat)
print(model_cat_cox)
print(model_lgb)
```



```

XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.5, device='cuda', early_stopping_roun
unds=None,
              enable_categorical=True, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.02, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=3, max_leaves=None,
              min_child_weight=80, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=2000, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.5, device='cuda', early_stopping_roun
unds=None,
              enable_categorical=True, eval_metric='cox-nloglik',
              feature_types=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=0.02, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=3,
              max_leaves=None, min_child_weight=80, missing=nan,
              monotone_constraints=None, multi_strategy=None, n_estimators=2000,
              n_jobs=None, num_parallel_tree=None, objective='survival:cox', ...)
<catboost.core.CatBoostRegressor object at 0x7abdd279c820>
<catboost.core.CatBoostRegressor object at 0x7abde0123820>
LGBMRegressor(colsample_bytree=0.4, device='gpu', learning_rate=0.02,
              max_depth=3, n_estimators=2500, objective='regression',
              verbose=-1)

```

In [30]:

```
import joblib
import lightgbm as lgb
import json
import os

# Save XGBoost Models
joblib.dump(model_xgb, "/kaggle/working/xgb_model.pkl")
joblib.dump(model_xgb_cox, "/kaggle/working/xgb_cox_model.pkl")
print("✅ XGBoost models saved!")

# Save CatBoost Models
model_cat.save_model("/kaggle/working/cat_model.cbm")
model_cat_cox.save_model("/kaggle/working/cat_cox_model.cbm")
print("✅ CatBoost models saved!")

# Save LightGBM Model
model_lgb.booster_.save_model("/kaggle/working/lgb_model.txt")
print("✅ LightGBM model saved!")

# Save Feature List
with open("/kaggle/working/features.json", "w") as f:
    json.dump(FEATURES, f)
print("✅ Feature list saved!")

# Verify saved files
print("📁 Saved files:", os.listdir("/kaggle/working/"))
```

✅ XGBoost models saved!

✅ CatBoost models saved!

✅ LightGBM model saved!

✅ Feature list saved!

📁 Saved files: ['.virtual_documents', 'cat_model.cbm', 'features.json', 'submission.csv', 'catboost_info', 'xgb_model.pkl', 'xgb_cox_model.pkl', 'lgb_model.txt', 'cat_cox_model.cbm']

In [31]:

```
!zip -r models.zip /kaggle/working/
```

```
adding: kaggle/working/ (stored 0%)
adding: kaggle/working/.virtual_documents/ (stored 0%)
adding: kaggle/working/cat_model.cbm (deflated 69%)
adding: kaggle/working/features.json (deflated 59%)
adding: kaggle/working/submission.csv (deflated 20%)
adding: kaggle/working/catboost_info/ (stored 0%)
adding: kaggle/working/catboost_info/learn/ (stored 0%)
adding: kaggle/working/catboost_info/learn/events.out.tfevents (d
eflated 74%)
adding: kaggle/working/catboost_info/learn_error.tsv (deflated 5
7%)
adding: kaggle/working/catboost_info/time_left.tsv (deflated 51%)
adding: kaggle/working/catboost_info/test/ (stored 0%)
adding: kaggle/working/catboost_info/test/events.out.tfevents (de
flated 75%)
adding: kaggle/working/catboost_info/test_error.tsv (deflated 6
4%)
adding: kaggle/working/catboost_info/tmp/ (stored 0%)
adding: kaggle/working/catboost_info/catboost_training.json (defl
ated 76%)
adding: kaggle/working/xgb_model.pkl (deflated 82%)
adding: kaggle/working/xgb_cox_model.pkl (deflated 81%)
adding: kaggle/working/lgb_model.txt (deflated 72%)
adding: kaggle/working/cat_cox_model.cbm (deflated 65%)
```

In [32]:

```

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
import seaborn as sns
import matplotlib.pyplot as plt

# Define a threshold for binary classification
threshold = 0.5 # Change to np.median(oof_xgb) if needed

# Convert survival scores to binary predictions
y_true = train["efs"] # Actual event labels

# Convert model outputs into binary labels using the threshold
y_pred_xgb = (oof_xgb > threshold).astype(int)
y_pred_cat = (oof_cat > threshold).astype(int)
y_pred_lgb = (oof_lgb > threshold).astype(int)
y_pred_ensemble = (rankdata(oof_xgb) + rankdata(oof_cat) + rankdata(oof_lgb) > np.median(rankdata(oof_xgb) + rankdata(oof_cat) + rankdata(oof_lgb))).astype(int)

# Function to evaluate model performance
def evaluate_model(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)

    # Display Confusion Matrix
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No Event", "Event"], yticklabels=["No Event", "Event"])
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(f"Confusion Matrix - {model_name}")
    plt.show()

    print(f" ♦ Model: {model_name}")
    print(f" ✓ Accuracy: {acc:.4f}")
    print(f" ✓ Precision: {prec:.4f}")
    print(f" ✓ Recall: {rec:.4f}")
    print(f" ✓ F1 Score: {f1:.4f}")
    print("- * 40)

```

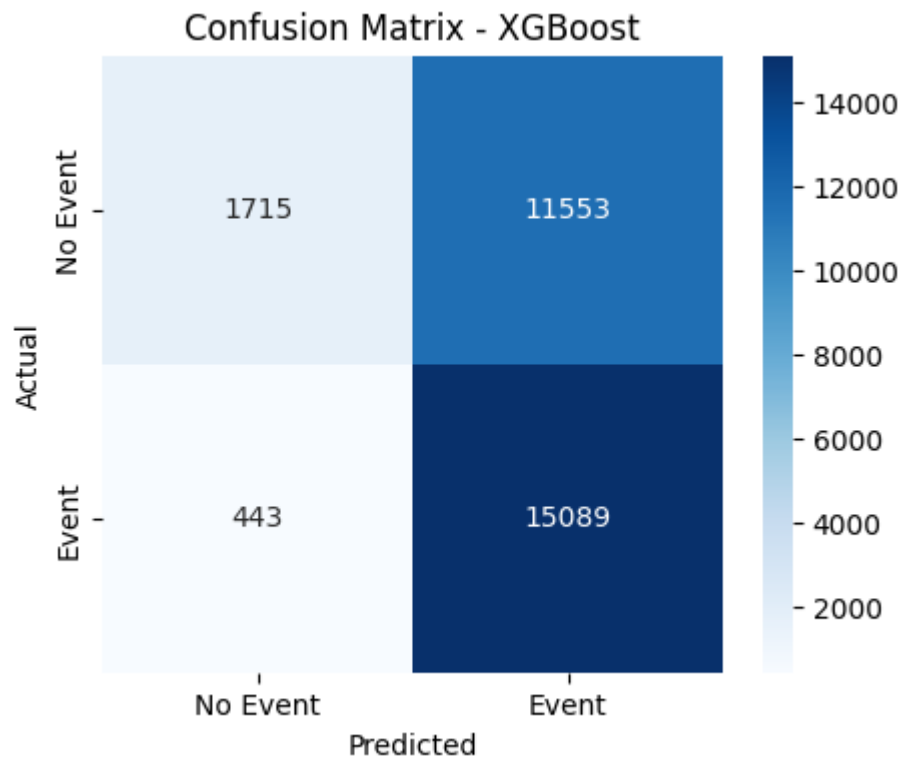
```
# Evaluate all models
```

```
evaluate_model(y_true, y_pred_xgb, "XGBoost")
```


```
evaluate_model(y_true, y_pred_cat, "CatBoost")
```

```
evaluate_model(y_true, y_pred_lgb, "LightGBM")
```

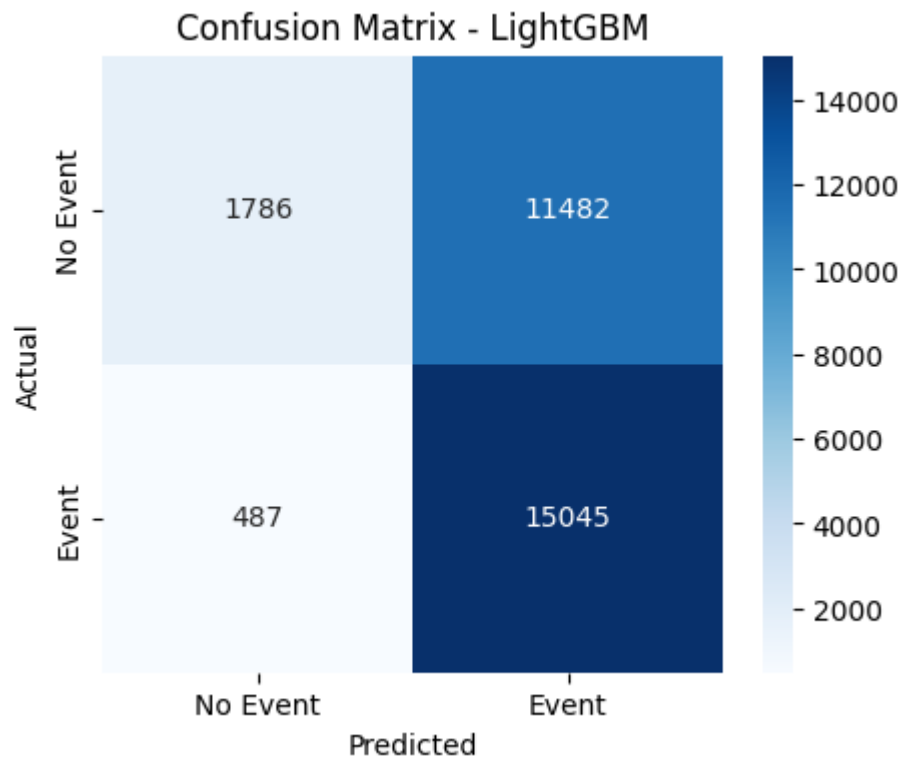
```
evaluate_model(y_true, y_pred_ensemble, "Ensemble (XGB + CAT + LGB)")
```



- ◆ Model: XGBoost
 - ✓ Accuracy: 0.5835
 - ✓ Precision: 0.5664
 - ✓ Recall: 0.9715
 - ✓ F1 Score: 0.7156
-




- ◆ Model: CatBoost
 - ✓ Accuracy: 0.5842
 - ✓ Precision: 0.5674
 - ✓ Recall: 0.9645
 - ✓ F1 Score: 0.7145
-



- ◆ Model: LightGBM
 - ✓ Accuracy: 0.5844
 - ✓ Precision: 0.5672
 - ✓ Recall: 0.9686
 - ✓ F1 Score: 0.7154
-

Confusion Matrix - Ensemble (XGB + CAT + LGB)



◆ Model: Ensemble (XGB + CAT + LGB)

✓ Accuracy: 0.6368

✓ Precision: 0.6761

✓ Recall: 0.6268

✓ F1 Score: 0.6505

In [33]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
import seaborn as sns
import matplotlib.pyplot as plt

# Separate categorical and numerical features
categorical_cols = train[FEATURES].select_dtypes(include=['category', 'object']).columns
numerical_cols = train[FEATURES].select_dtypes(exclude=['category', 'object']).columns

# Fill missing values
train_filled = train[FEATURES].copy()
train_filled[categorical_cols] = train_filled[categorical_cols].fillna(train_filled[categorical_cols].mode().iloc[0])
train_filled[numerical_cols] = train_filled[numerical_cols].fillna(train_filled[numerical_cols].median())

# Train Random Forest Classifier
model_rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=5,
    random_state=42,
    n_jobs=-1
)

# Train the model
model_rf.fit(train_filled, train["efs"])


# Predict probabilities and convert to binary classification
oof_rf = model_rf.predict_proba(train_filled)[: , 1] # Probability of class 1 (event)
y_pred_rf = (oof_rf > 0.5).astype(int) # Convert to binary classification

# Function to evaluate model performance
def evaluate_model(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
```

```
# Display Confusion Matrix
plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No
Event", "Event"], yticklabels=["No Event", "Event"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Confusion Matrix - {model_name}")
plt.show()

print(f" ♦ Model: {model_name}")
print(f" ✓ Accuracy: {acc:.4f}")
print(f" ✓ Precision: {prec:.4f}")
print(f" ✓ Recall: {rec:.4f}")
print(f" ✓ F1 Score: {f1:.4f}")
print("-" * 40)

# Evaluate Random Forest model
evaluate_model(train["efs"], y_pred_rf, "Random Forest")
```



◆ Model: Random Forest

✓ Accuracy: 0.6577

✓ Precision: 0.6354

✓ Recall: 0.8571

✓ F1 Score: 0.7298

In [34]:

```
import matplotlib.pyplot as plt
import numpy as np

# Model names
models = ["XGBoost", "CatBoost", "LightGBM", "Ensemble", "Random Forest"]

# Metrics for each model
accuracy = [0.5835, 0.5842, 0.5845, 0.6369, 0.6577]
precision = [0.5664, 0.5674, 0.5672, 0.6762, 0.6354]
recall = [0.9715, 0.9645, 0.9686, 0.6269, 0.8571]
f1_score = [0.7156, 0.7145, 0.7154, 0.6506, 0.7298]

# Set bar positions
x = np.arange(len(models))
width = 0.2 # Bar width

plt.figure(figsize=(10, 6))

# Plot bars for each metric
bars1 = plt.bar(x - 1.5*width, accuracy, width, label='Accuracy', color='blue', alpha=0.7)
bars2 = plt.bar(x - 0.5*width, precision, width, label='Precision', color='green', alpha=0.7)
bars3 = plt.bar(x + 0.5*width, recall, width, label='Recall', color='red', alpha=0.7)
bars4 = plt.bar(x + 1.5*width, f1_score, width, label='F1 Score', color='purple', alpha=0.7)

# Function to add labels on top of bars
def add_labels(bars):
    for bar in bars:
        height = bar.get_height()
        plt.text(bar.get_x() + bar.get_width()/2, height + 0.02, f'{height:.3f}', ha='center', fontsize=10)

# Add values to bars
add_labels(bars1)
add_labels(bars2)
add_labels(bars3)
add_labels(bars4)

# Labels and titles
```

```
plt.xlabel("Models")
plt.ylabel("Metric Scores")
plt.title("Model Performance Comparison")
plt.xticks(x, models)
plt.ylim(0, 1.1) # Extend y-axis slightly to fit labels

# Add legend
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Show the plot
plt.show()
```

